

**ECS7026P**

**Neural Networks and Deep Learning**

Coursework Assignment

**NITISH KRISHNA SADHU**

220283937

EC22098

## Task – 1:

The following code downloads the data and applies the defined transformation.

```
# Transformation for the CIFAR-10 dataset
transformation = transforms.Compose(
    [transforms.RandomHorizontalFlip(), # Randomly flipping the image
     transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

# Loading the CIFAR-10 training dataset
trainSet = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=transformation)

# Loading the CIFAR-10 test dataset
testSet = torchvision.datasets.CIFAR10(root='./data", train=False, download=True, transform=transformation)
```

In the following code:

- The *batch\_size* is defined as 256.
- The next line of the code converts the *trainSet* into a *trainLoader* which is in the dataloader format.
- The final line of the code converts the *testSet* into a *testLoader* which is in the dataloader format.

```
# Defining the batch size
batch_size = 256

# Loading the datasets
# Training dataset
trainLoader = torch.utils.data.DataLoader(trainSet, batch_size=batch_size, shuffle=True, num_workers=2, pin_memory=2)

# Test dataset
testLoader = torch.utils.data.DataLoader(testSet, batch_size=batch_size, shuffle=True, num_workers=2, pin_memory=2)
```

## Task-2:

The model:

- Consists of four **BackboneBlocks**.
- Each **BackboneBlock** consists of two parts:
  - The first part consists of the code that produces the vector **a**.
  - The second part consists of the convolution layers.
- In the forward function, the vector **a** is computed first.
  - The input first goes into an average pooling layer.
  - The output then goes into a Linear layer which outputs **k** features.
  - The ReLU activation is applied to the output.
- The next step is to apply the convolution layers to the input images and multiply the output of each convolution layer with the individual elements of the vector.
- The next step is to apply the batch normalization and max pooling which helps in improving the accuracy.
- The output from the four **BackboneBlocks** is comes out, it goes into the classifier part of the model which is an MLP layer, with two layers.

- The **ReLU** activation is given between the two layers.
- A dropout with a  $p=0.25$  is also given between the two layers.

### Task-3:

- The loss used is CrossEntropyLoss.
- The optimizer used is Adam with a learning rate of 0.001.


### Task-4:

- Epochs used: 20
- The learning rate used: 0.001.
  - Tested with the learning rates 0.1, 0.01, 0.001, 0.0001.
  - The accuracy was lower for 0.1 and 0.01 than 0.001.
  - The time taken to train 0.0001 is high.
  - The Learning rate Of 0.001 yielded the best accuracy.
- The batch size used: 256.
  - Tested with 4, 16, 128, 256.
  - The batch size of 256 yielded the best results.
- The k used is 16.
  - Tested with  $k = 8, 16, 32$ .
  - $K=16$  yielded the better results among the three.

### # Graph of loss, training accuracy and test accuracy.

```
[61] n_epochs = 20    # Defining the number of epochs
      lr = 0.001      # Defining the learning rate
      train_model(model, trainLoader, testLoader, n_epochs, lr, device)  # Training the model
```

loss 0.680, train acc 0.857, test acc 0.823  
3616.0 examples/sec on cuda



### Task-5:

- The accuracy initially obtained was around 30 percent.
- Has tested with 3, 4, 5 **BackboneBlocks**, 4 blocks gave the best of both world, i.e., training reasonably fast and providing the best accuracy.
- The learning rate used: 0.001.
  - Tested with the learning rates 0.1, 0.01, 0.001, 0.0001.
  - The accuracy was lower for 0.1 and 0.01 than 0.001.
  - The time taken to train 0.0001 is high.
  - The Learning rate Of 0.001 yielded the best accuracy.
- The batch size used: 256.
  - Tested with 4, 16, 128, 256.
  - The batch size of 256 yielded the best results.
- The k used is 16.
  - Tested with k = 8, 16 ,32.
  - K=16 yielded the better results among the three.
- The addition of Batch normalization and Max pooling before emitting the output from **BackboneBlock** drastically improved the accuracy.
- The classifier has two linear layers, using the ReLU activation and dropout with p=0.3, improved the accuracy from around 70% to around 82.3%.

### References:

- Used my\_utils.py file with some modifications to write the training script and plot the graphs.